

Securing Apache Kafka

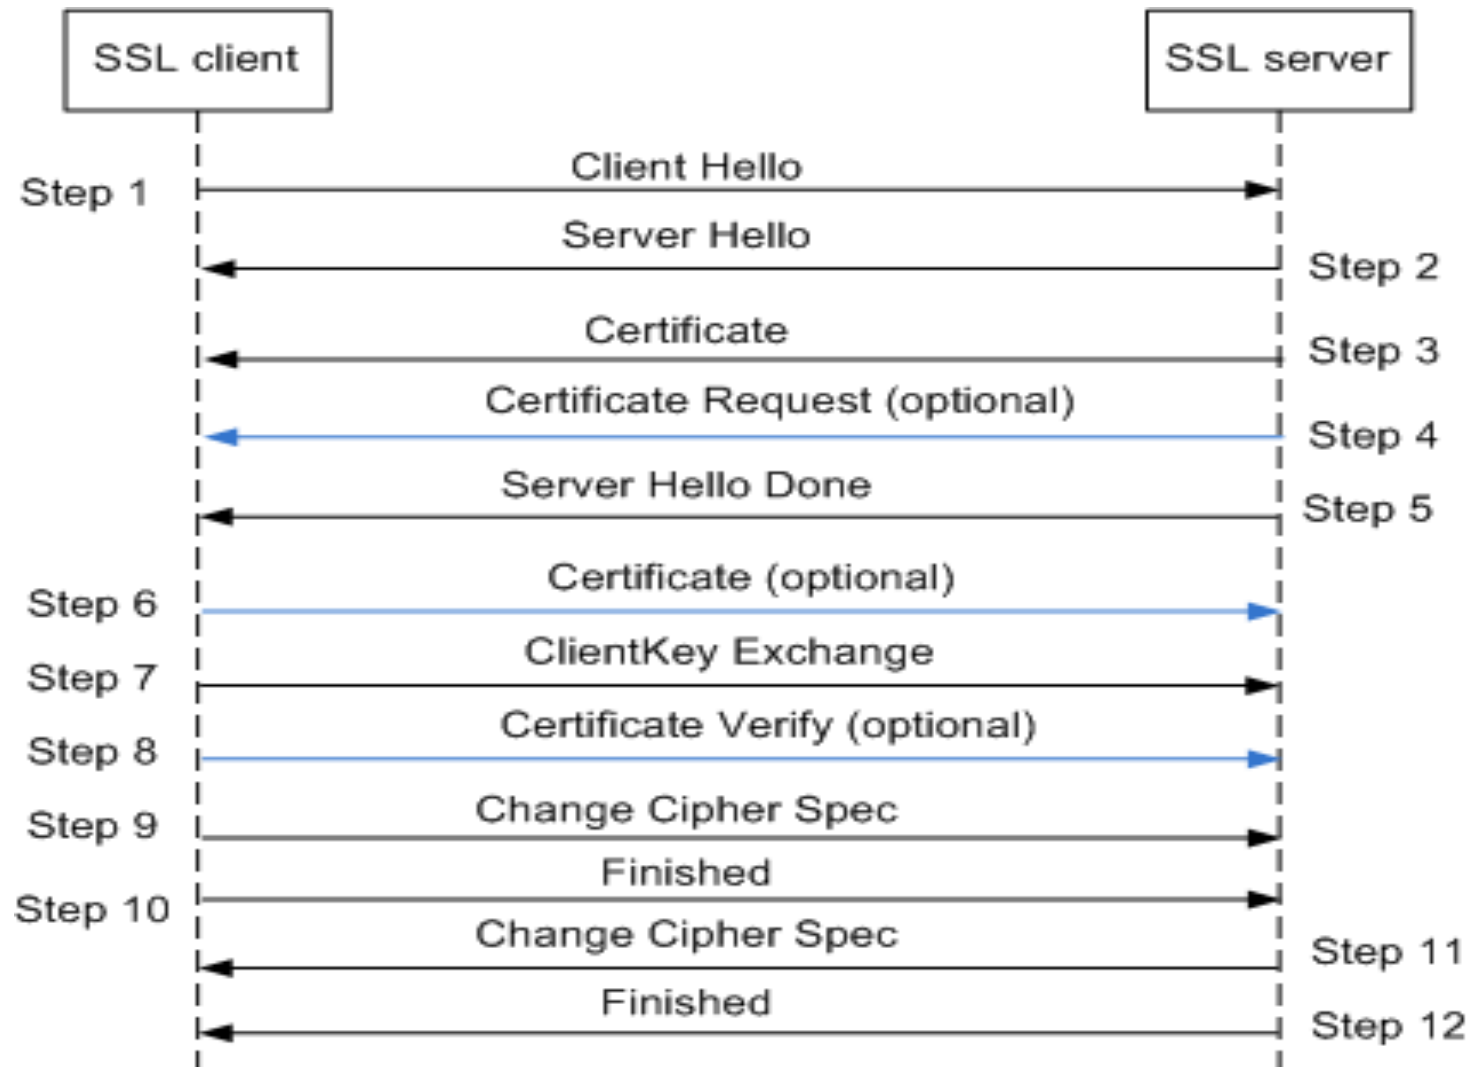
- Kafka and security overview
- Authentication
 - Identify the principal (user) associated with a connection
- Authorization
 - What permission a principal has
- Secure Zookeeper

- Support since 0.9.0
- Wire encryption btw client and broker
 - For cross data center mirroring
- Access control on resources such as topics
 - Enable sharing Kafka clusters

- Broker support multiple ports
 - plain text (no wire encryption/authentication)
 - SSL (for wire encryption/authentication)
 - SASL (for Kerberos authentication)
 - SSL + SASL (SSL for wire encryption, SASL for authentication)
- Clients choose which port to use
 - need to provide required credentials through configs

- 1-way authentication
 - Secure wire transfer through encryption
- 2-way authentication
 - Broker knows the identity of client
- Easy to get started
 - Just involve client and server

SSL handshake



- Data encrypted with agreed upon cipher suite
 - Encryption overhead
 - Losing zero-copy transfer in consumer

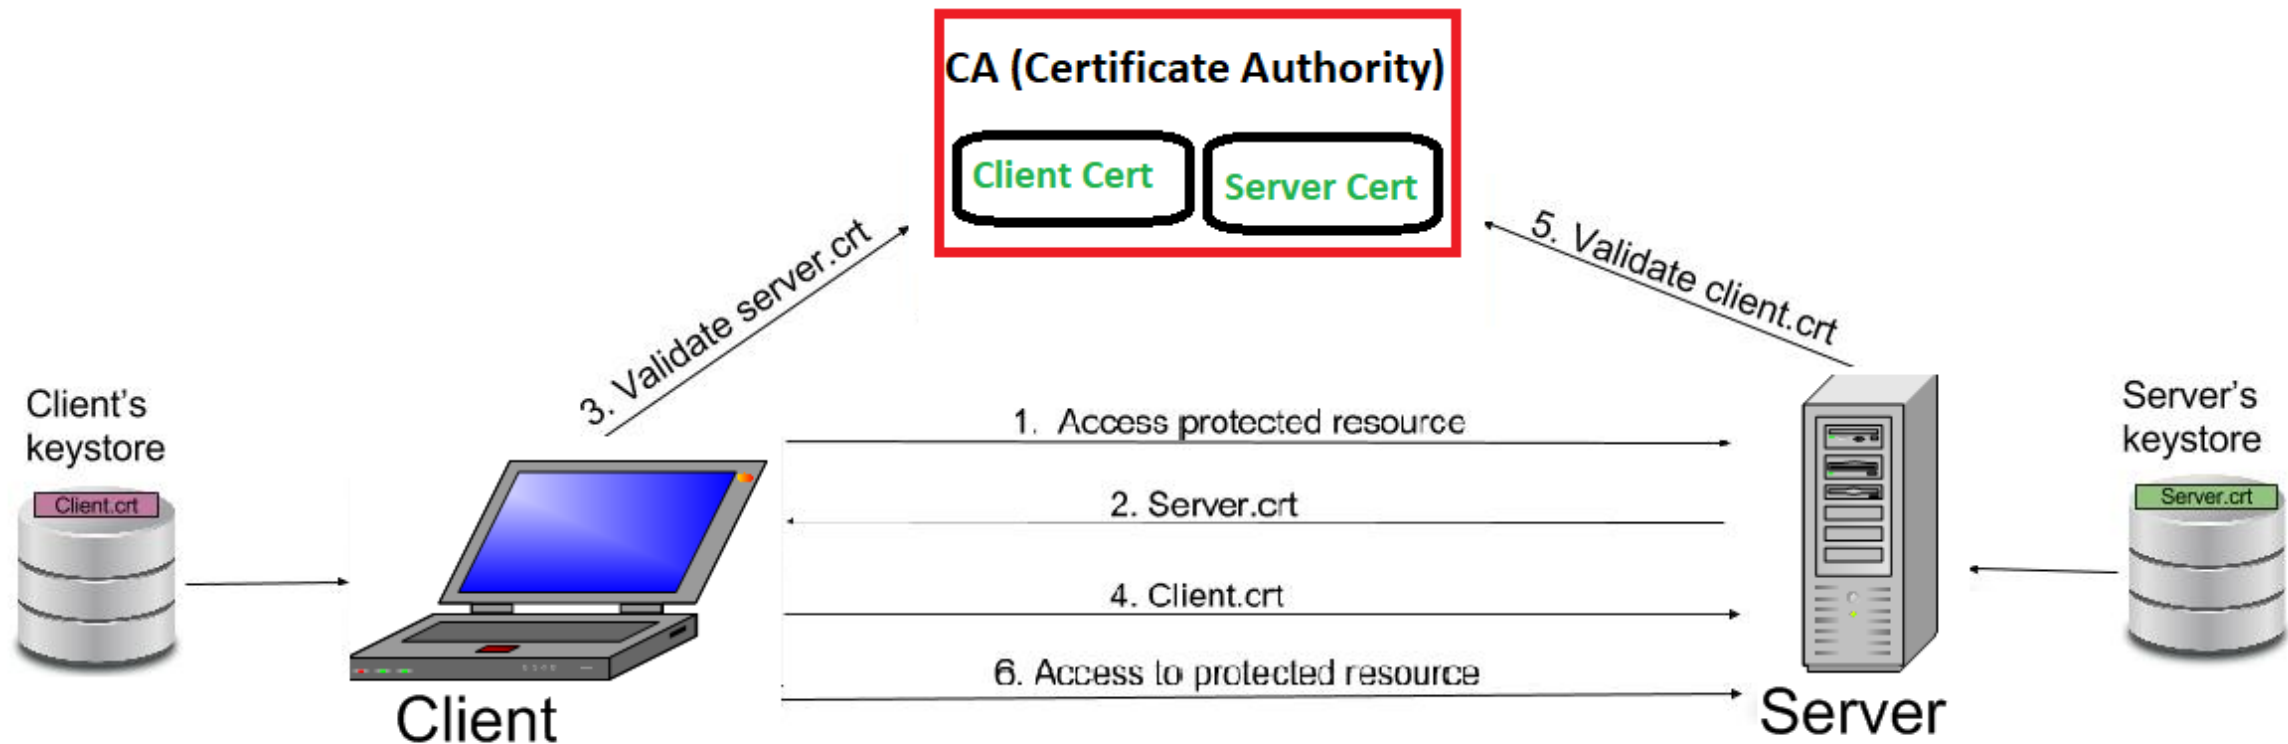
- **r3.xlarge**
 - 4 core, 30GB ram, 80GB ssd, moderate network (~90MB/s)

	Throughput(MB/S)	CPU on client	CPU on broker
Producer(plaintext)	83	12%	30%
Producer (SSL)	69	28%	48%
Consumer (plaintext)	83	8%	2%
Consumer (SSL)	69	27%	24%

- **Most overhead from encryption**

Preparing SSL

1. Generate certificate (X509) in broker key store
2. Generate certificate authority (CA) for signing
3. Sign broker certificate with CA
4. Import signed certificate and CA to broker key store
5. Import CA to client trust store
6. 2-way authentication: generate client certificate in a similar way



- No client code change; just configuration change.

Client/Broker

```
ssl.keystore.location =  
/var/private/ssl/kafka.server.keystore.jks  
ssl.keystore.password = test1234  
ssl.key.password = test1234  
ssl.truststore.location =  
/var/private/ssl/kafka.server.truststore.jks  
ssl.truststore.password = test1234
```

Broker

```
listeners = SSL://host.name:port  
security.inter.broker.protocol = SSL  
ssl.client.auth = required
```

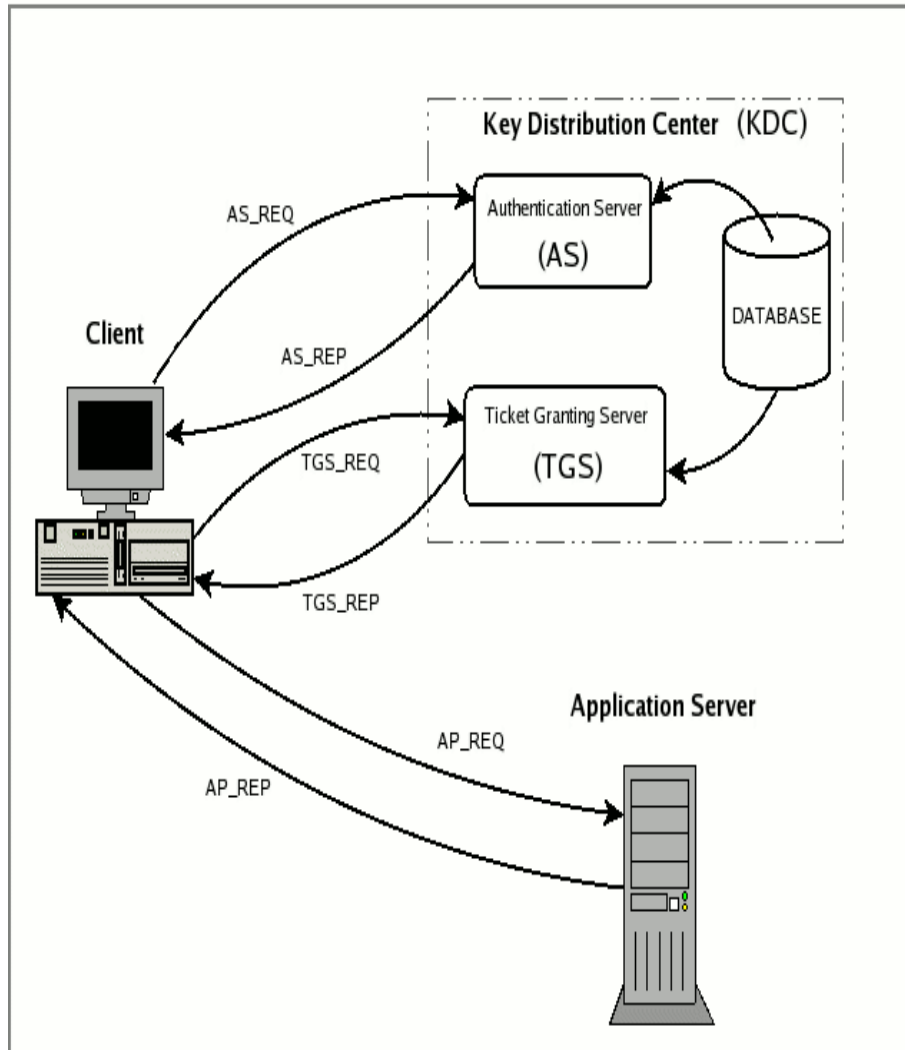
Client

```
security.protocol = SSL
```

- Simple Authentication and Security Layer
 - Challenge/response protocols
 - Server issues challenge and client sends response
 - Continue until server is satisfied
- Different mechanisms
 - Plain: cleartext username/password
 - Digest MD5
 - GSSAPI: Kerberos
- Kafka 0.9.0 only supports Kerberos

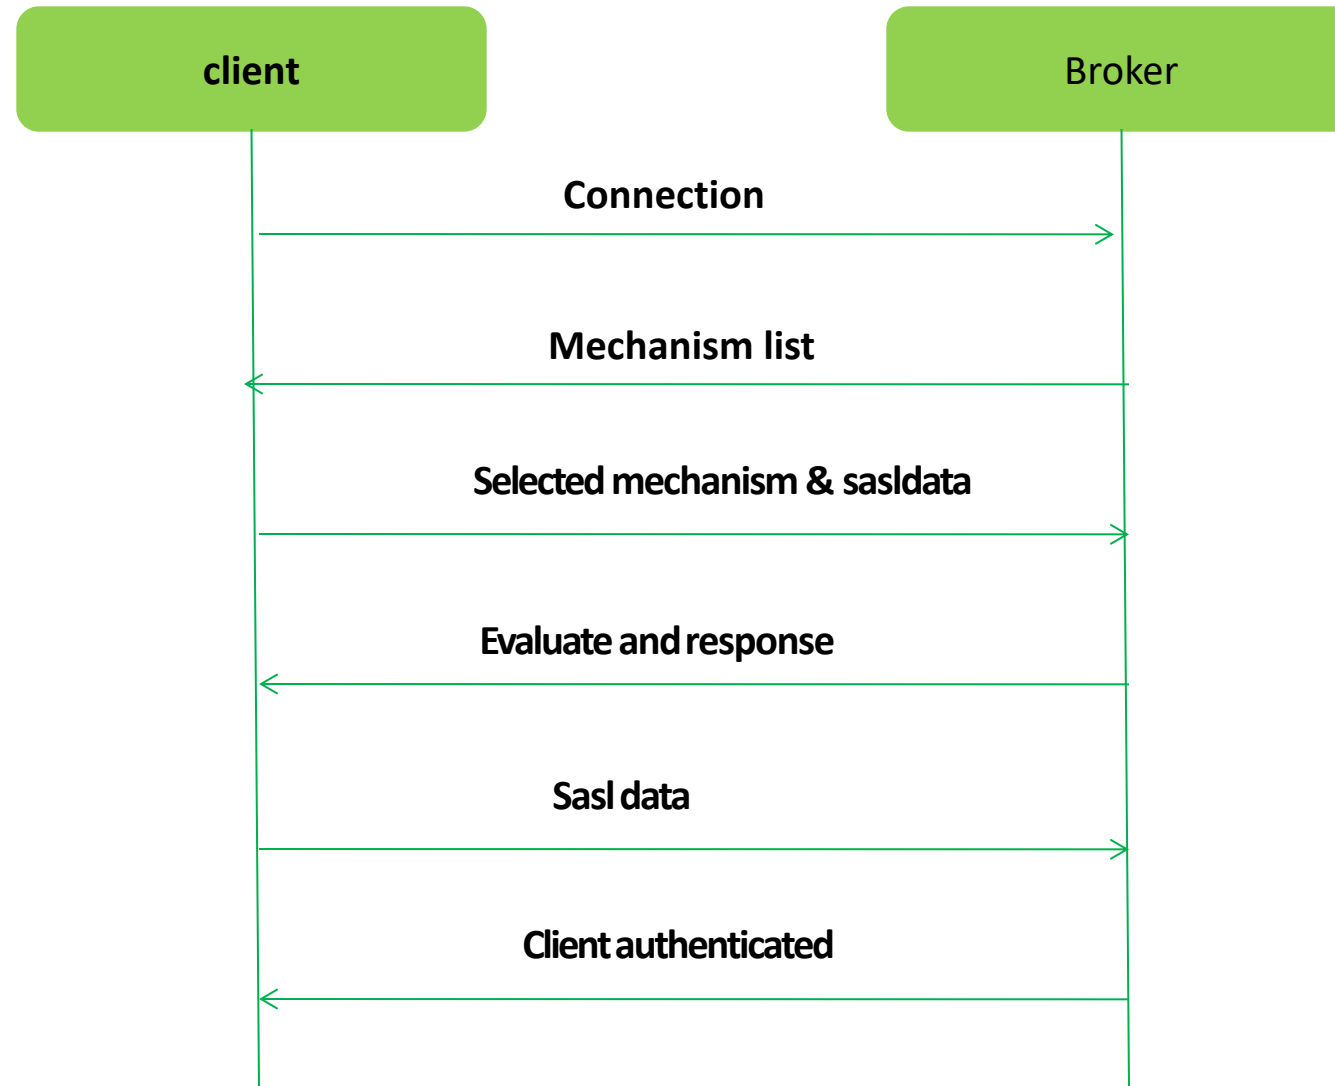
- Secure single sign-on
 - An organization may provide multiple services
 - User just remember a single Kerberos password to use all services
- More convenient when there are many users
- Need Key Distribution Center (KDC)
 - Each service/user need a Kerberos principal in KDC

How Kerberos Works



- Create service and client principal in KDC
- Client authenticate with AS on startup
- Client obtain service ticket from TGS
- Client authenticate with service using service ticket

SASL handshake



- SASL_PLAINTEXT
 - No wire encryption
- SASL_SSL
 - Wire encryption over SSL

- Create Kafka service principal in KDC
- Create a keytab for Kafka principal
 - Keytab includes principal and encrypted Kerberos password
 - Allow authentication w/o typing password
- Create an application principal for client KDC
- Create a keytab for application principal

Configuring Kerberos

No client code change; just configuration change

Broker JAAS file

```
KafkaServer {  
    com.sun.security.auth.module.  
    Krb5LoginModule required  
    useKeyTab=true  
    storeKey=true  
    keyTab="/etc/security/keyt  
    abs/kafka_server.keytab"  
    principal="kafka/kafka1.ho  
    stname.com@EXAMPLE.COM";  
};
```

Client JAAS file

```
KafkaClient {  
    com.sun.security.auth.module.  
    Krb5LoginModule required  
    useKeyTab=true  
    storeKey=true  
    keyTab="/etc/security/keyt  
    abs/kafka_client.keytab"  
    principal="kafka-client-  
    1@EXAMPLE.COM";  
};
```

Broker JVM

```
Djava.security.auth.lo  
gin.config=/etc/kafka/  
kafka_server_jaas.conf
```

ClientJVM

```
-  
Djava.security.auth.log  
in.config=/etc/kafka/  
kafka_client_jaas.conf
```

Broker config

```
security.inter.broker.protocol=  
SASL_PLAINTEXT(SASL_SSL)  
sasl.kerberos.service.name=kafka
```

Clientconfig


```
security.protocol=SA  
SL_PLAINTEXT(SASL_SSL)  
sasl.kerberos.servic  
e.name=kafka
```

- Kerberos principal
 - Primary[/Instance]@REALM
 - [kafka/kafka1.hostname.com@EXAMPLE.COM](#)
 - [kafka-client-1@EXAMPLE.COM](#)
- Primary extracted as the default principal name
- Can customize principal name through
sasl.kerberos.principal.to.local.rules

- Authentication (SSL or SASL) happens once during socket connection
 - No re-authentication
- If a certificate needs to be revoked, use authorization to remove permission

- Control which permission each authenticated principal has
- Pluggable with a default implementation

Alice is Allowed to Read from topic T1 from Host1



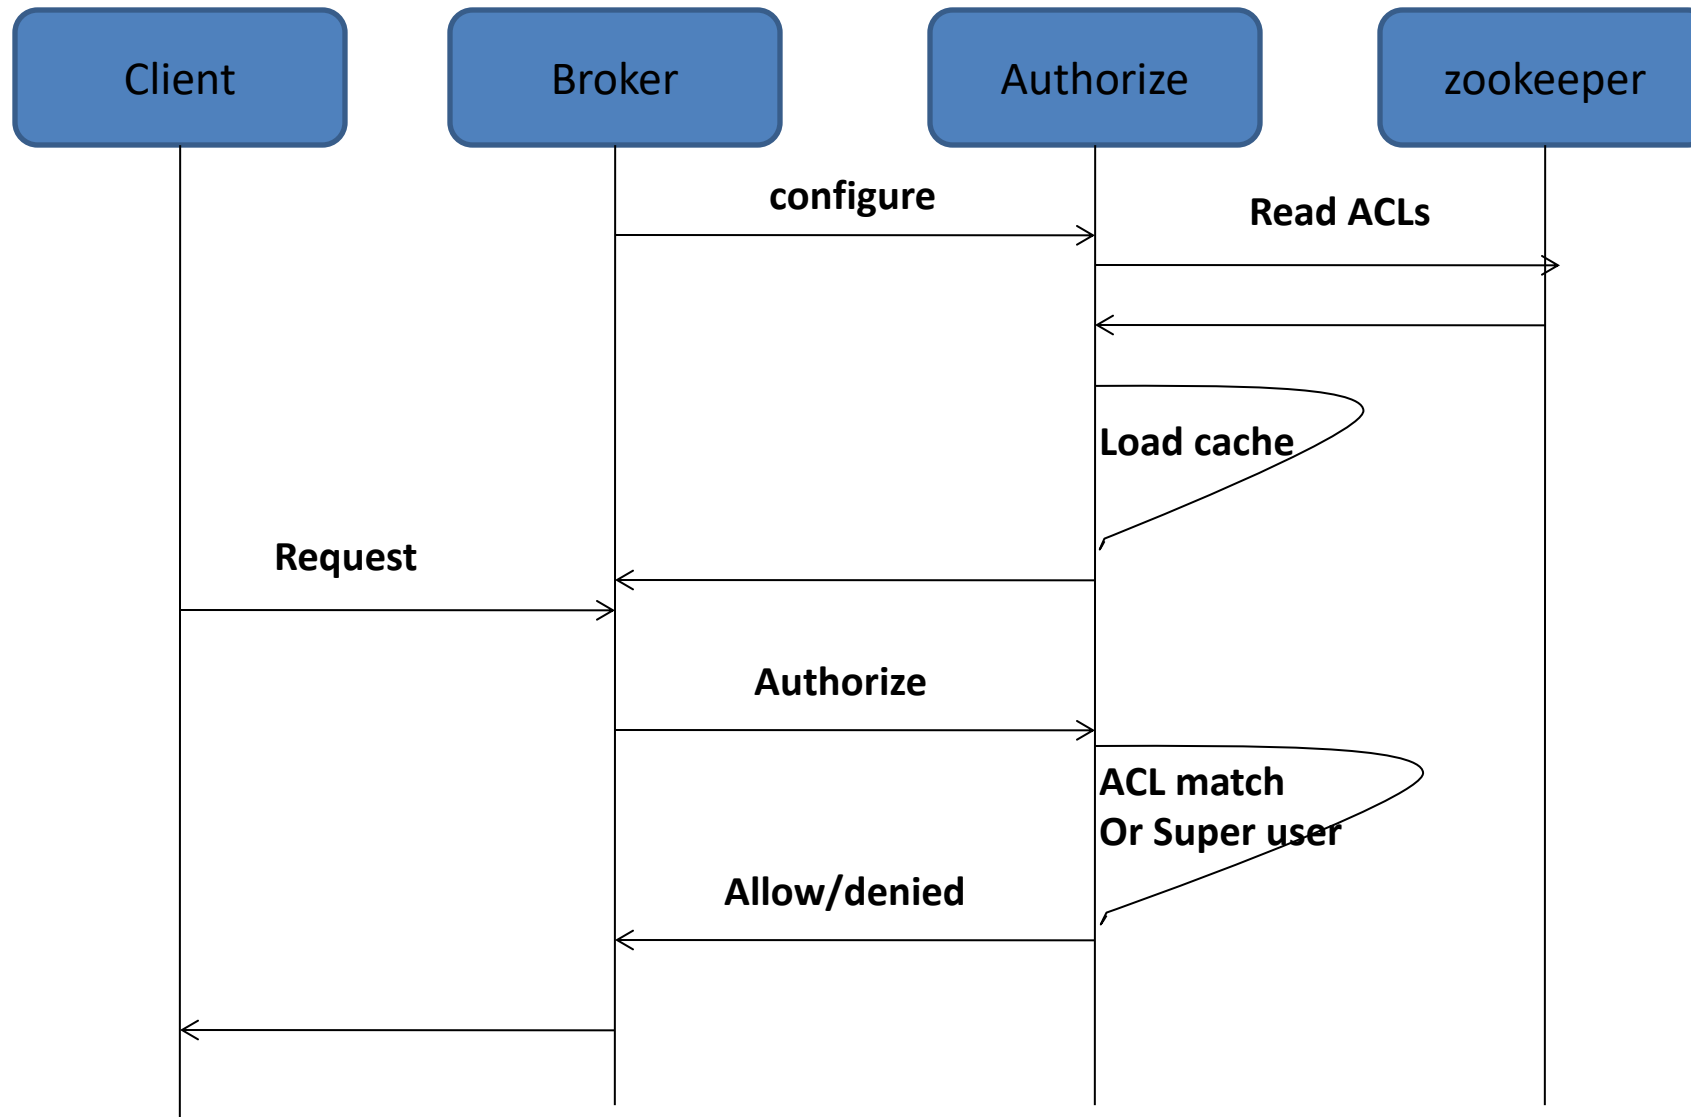
- Operations
 - Read, Write, Create, Describe, ClusterAction, All
- Resources
 - Topic, Cluster and ConsumerGroup

Operations	Resources
Read, write, Describe (Read, Write implies Describe)	Topic
Read	Consumer Group
Create, ClusterAction(communication between controller and brokers)	Cluster

- Out of box authorizer implementation.
- CLI tool for adding/removing acls
- ACLs stored in zookeeper and propagated to brokers asynchronously
- ACL cache in broker for better performance

Authorizer Flow

Tos



- `authorizer.class.name=kafka.security.auth.SimpleAclAuthorizer`
- Make Kafka principal super users
 - Or grant ClusterAction and Read all topics to Kafka principal

- Producer

- Grant Write on topic, Create on cluster (auto creation)
- Or use --producer option in CLI

```
bin/kafka-acls --authorizer-properties zookeeper.connect=localhost:2181 \  
--add --allow-principal User:Bob --producer --topic t1
```

- Consumer

- Grant Read on topic, Read on consumer group
- Or use --consumer option in CLI

```
bin/kafka-acls --authorizer-properties zookeeper.connect=localhost:2181 \  
--add --allow-principal User:Bob --consumer --topic t1 --group group1
```

- Zookeeper stores
 - critical Kafka metadata
 - ACLs
- Need to prevent untrusted users from modifying

- ZK supports authentication through SASL
 - Kerberos or Digest MD5
- Set `zookeeper.set.acl` to true on every broker
- Configure ZK user through JAAS config file
 - Each ZK path writable by creator, readable by all

- Configure brokers with multiple ports
 - `listeners=PLAINTEXT://host.name:port,SSL://host.name:port`
- Gradually migrate clients to secure port
- When done
 - Turn off PLAINTEXT port on brokers

Lab : - Securing Kafka